

Stage 2 — per-group metrics (fine taxonomy, held-out, oracle n=89)

group	n	hits	precision	recall (hit)	F1	accuracy ¹ (1-vs-rest)
black_bars	28	14	44%	50%	0.47	64%
overlay	27	10	33%	37%	0.35	58%
empty_space	16	3	33%	19%	0.24	79%
disordered_layout	13	3	17%	23%	0.19	72%
⚠ distortion	3	0	—	0%	0.00	97%
⚠ orientation	2	0	—	0%	0.00	98%
GLOBAL (micro)	89	30	34%	34%	0.34	34%

Precision = of those the model called X, how many were X (— = never predicted X: 0/0 undefined). Recall (hit rate) = of the real X, how many were caught. Accuracy¹ = (hit+correct reject)/total (one-vs-rest).

¹ inflated by true-negatives on rare classes — ALWAYS read with the support (n). ⚠ = support < 5 (unstable; read the raw fraction hits/n). GLOBAL row = micro (precision = recall = F1 = accuracy = 34%, multiclass identity); macro-F1 = 0.21 (in title).