

Stage 2 — per-group metrics (fine taxonomy, held-out, oracle n=89)

group	n	hits	precision	recall (hit)	F1	accuracy ¹ (1-vs-rest)
black_bars	28	16	52%	57%	0.54	70%
overlay	27	9	35%	33%	0.34	61%
empty_space	16	4	31%	25%	0.28	76%
disordered_layout	13	3	18%	23%	0.20	73%
⚠ distortion	3	2	100%	67%	0.80	99%
⚠ orientation	2	0	—	0%	0.00	98%
GLOBAL (micro)	89	34	38%	38%	0.38	38%

Precision = of those the model called X, how many were X (— = never predicted X: 0/0 undefined). Recall (hit rate) = of the real X, how many were caught. Accuracy¹ = (hit+correct reject)/total (one-vs-rest).

¹ inflated by true-negatives on rare classes — ALWAYS read with the support (n). ⚠ = support < 5 (unstable; read the raw fraction hits/n). GLOBAL row = micro (precision = recall = F1 = accuracy = 38%, multiclass identity); macro-F1 = 0.36 (in title).